

Easy Assignment – 1 Disha (2022CS11118)

Part 2: Enhanced Shell

Objective: Implement a secure login system for the xv6 shell, requiring a username and password. The system allows three login attempts, after which it locks out the user.

Implementation: (init.c)

1. Username Prompt & Validation:

- The user is prompted to enter a username.
- The input is checked against the predefined USERNAME. If incorrect, the user is asked to try again until the correct username is entered.

```
while (strcmp(input_username, USERNAME) != 0) {  
    printf(1, "Invalid username. Try again.\n");  
    printf(1, "Enter Username: ");  
    gets(input_username, sizeof(input_username));  
    input_username[strlen(input_username) - 1] = '\0';  
}
```

Figure 1: init.c

2. Password Prompt & Validation:

- After a correct username, the user is prompted to enter a password.
- The input password is compared to the predefined PASSWORD.
- If incorrect, the user has three attempts to enter the correct password.

```
while (attempts < MAX_ATTEMPTS) {  
    printf(1, "Enter Password: ");  
    gets(input_password, sizeof(input_password));  
    input_password[strlen(input_password) - 1] = '\0';  
  
    if (strcmp(input_password, PASSWORD) == 0) {  
        printf(1, "Login successful\n");  
        return 1;  
    } else {  
        attempts++;  
        printf(1, "Incorrect password. Attempt %d of %d\n", attempts, MAX_ATTEMPTS);  
    }  
}  
printf(1, "Login failed. System locked.\n");  
return 0;
```

Figure 2: init.c

3. Attempts Limitation:

- If the user exceeds three incorrect password attempts, the system locks the login process.

```
if (login() == 0) {  
    exit();  
}
```

Figure 3: init.c

4. Success:

- Upon correct username and password entry, the user gains access, and the shell proceeds.

```
Enter Username: 2022CS11118
Enter Password: disha
Login successful
init: starting sh
$
```

Part 3: History

Objective: To implement an additional command history that needs to display a list of all the processes that have been executed until now

Implementation :

- `history_entry` structure is defined to store the required process details (PID, name, total memory).

```
5
6 #define HISTORY_SIZE 100
7
8 struct history_entry {
9     int pid;
10    char name[16];
11    uint total_mem;
12 };
13
14 struct history_entry history[HISTORY_SIZE];
15 int history_count = 0;
16
17 void
```

Figure 4: *proc.c*

- maintained an array of `history_entry` structures and a count to track how many processes are stored in the history

```
if (history_count < HISTORY_SIZE) {
    history[history_count].pid = curproc->pid;
    safestrcpy(history[history_count].name, curproc->name, sizeof(curproc->name));
    history[history_count].total_mem = curproc->sz;
    history_count++;
}
```

Figure 5: *proc.c* : `exit()`

```
int gethistory() {
    for(int i = 0; i < history_count; i++) {
        cprintf("%d %s %d\n", history[i].pid, history[i].name, history[i].total_mem);
    }
    return 0;
}
```

Figure 6: *proc.c* (helper function to print the history)

Part 4 : block and unblock system calls

Objective: need a way to keep track of which system calls are blocked for each process spawned by the shell.
Implementation: Used a global array to store blocked system calls for processes associated with the shell.

```
char name[16]; // Process name (debugging)
int blocked_syscalls[MAX_SYSCALLS]; // Array for blocking syscalls for this process
int blocked_child_syscalls[MAX_SYSCALLS];
};
```

Figure 7: proc.h : struct proc

```
int block(int syscall_id) {
    if ((syscall_id < 0 || syscall_id >= MAX_SYSCALLS) || is_restricted(syscall_id)) {
        return -1;
    }
    struct proc *curproc = myproc();
    curproc->blocked_child_syscalls[syscall_id] = 1;
    return 0;
}
```

Figure 8: proc.c : block.c

similarly 'int unblock(int syscall_id)' is implemented

```
static int restricted_syscalls[] = {1, 2}; //cannot block fork and exit

int is_restricted(int syscall_id) {
    int i;
    for (i = 0; i < sizeof(restricted_syscalls) / sizeof(restricted_syscalls[0]); i++)
        if (restricted_syscalls[i] == syscall_id) {
            return 1; //
        }
    return 0;
}
```

Figure 9: proc.c : some system calls cannot be blocked

```
// blocked syscalls for the child
for(int i = 0; i < MAX_SYSCALLS; i++){
    np->blocked_syscalls[i] = curproc->blocked_child_syscalls[i];
    np->blocked_child_syscalls[i] = curproc->blocked_child_syscalls[i];
}
```

Figure 10: proc.c : child processes inheriting blocked syscalls from the parent, and the syscalls for the child processes of that child are also blocked.

```

4 struct proc *curproc = myproc();
5
6 num = curproc->tf->eax;
7 if (!(curproc->pid == 2 || strncmp(curproc->name, "sh", sizeof("sh") - 1) == 0)) {
8     if (curproc->blocked_syscalls[num] == 1) {
9         cprintf("syscall %d is blocked\n", num);
10        return;
11    }
12 }
13
14 if (num > 0 && num < NELEM(syscalls) && syscalls[num]) {

```

Figure 11: syscall.c

system calls are not blocked for the current shell, they are blocked only for the child processes

part 5: chmod

Objective: Implement the chmod system call to modify file permissions using a 3-bit integer (read, write, execute) in the xv6 operating system.

```
int chmod(const char* file, int mode) {  
    struct inode *ip = namei((char*)file);  
  
    begin_op();  
  
    if ((ip = namei((char*)file)) == 0) {  
        end_op();  
        return -1;  
    }  
  
    ilock(ip);  
  
    ip->mode = mode & 0b111;  
  
    iupdate(ip);  
    iunlock(ip);  
    end_op();  
  
    return 0;  
}
```

Figure 12: proc.c

Implementation:

1. Define the chmod Function:

- The function takes two arguments: a file (filename) and mode (3-bit permission value).
- The file is located using the namei() function, which returns the corresponding inode.

2. Check File Existence:

- If the file does not exist, the function returns -1 and terminates.

```
// On-disk inode structure  
struct dinode {  
    short type;           // F  
    short major;         // M  
    short minor;         // M  
    short nlink;         // N  
    uint size;           // S  
    uint addrs[NDIRECT+1]; // A  
    uint mode;           // M  
    char padding[60];    // Pad  
};
```

Figure 13: fs.h

Created a member “mode” to store the mode in disk. Also padding was done.

```
short nlink;  
uint size;  
uint addrs[NDIRECT+1];  
  
uint mode;  
};
```

Figure 14: file.h – struct inode

Created a member “mode” to store the mode of the file in memory.

3. Lock the Inode:

- If the file is found, the inode is locked using `ilock()`. This ensures no other process can modify the inode while it's being updated.

```
void
ilock(struct inode *ip)
{
    struct buf *bp;
    struct dinode *dip;

    if(ip == 0 || ip->ref < 1)
        panic("ilock");

    acquiresleep(&ip->lock);

    if(ip->valid == 0){
        bp = bread(ip->dev, IBLOCK(ip->inum, sb));
        dip = (struct dinode*)bp->data + ip->inum%IPB;
        ip->type = dip->type;
        ip->major = dip->major;
        ip->minor = dip->minor;
        ip->nlink = dip->nlink;
        ip->size = dip->size;
        memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
        ip->mode = dip->mode; //disk to memory
        brelse(bp);
        ip->valid = 1;
    }
```

Figure 15: *fs.c*

4. Update Permissions:

- The permission bits are masked with `0b111` to ensure only the last three bits are set (read, write, and execute permissions).
- The mode is then updated in the inode.

```
void
iupdate(struct inode *ip)
{
    struct buf *bp;
    struct dinode *dip;

    bp = bread(ip->dev, IBLOCK(ip->inum, sb));
    dip = (struct dinode*)bp->data + ip->inum%IPB;
    dip->type = ip->type;
    dip->major = ip->major;
    dip->minor = ip->minor;
    dip->nlink = ip->nlink;
    dip->size = ip->size;
    memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
    dip->mode = ip->mode; //memory to disk
    log_write(bp);
    bwrite(bp);
    brelse(bp);
}
```

Figure 16: *fs.c*

5. Write Back to Disk:

- After modifying the mode, the inode is updated on disk using `iupdate()`.

```
uint
ialloc(ushort type)
{
    uint inum = freeinode++;
    struct dinode din;

    bzero(&din, sizeof(din));
    din.type = xshort(type);
    din.nlink = xshort(1);
    din.size = xint(0);
    din.mode = 0x7;
    winode(inum, &din);
    return inum;
}
```

Figure 18: `mkfs.c`

```
struct inode*
ialloc(uint dev, short type)
{
    int inum;
    struct buf *bp;
    struct dinode *dip;

    for(inum = 1; inum < sb.ninodes; inum++){
        bp = bread(dev, IBLOCK(inum, sb));
        dip = (struct dinode*)bp->data + inum%IPB;
        if(dip->type == 0){ // a free inode
            memset(dip, 0, sizeof(*dip));
            dip->type = type;
            dip->mode = 0x7;
            log_write(bp); // mark it allocated on the disk
            brelse(bp);
            return iget(dev, inum);
        }
        brelse(bp);
    }
    panic("ialloc: no inodes");
}
```

Figure 17: `fs.c`

```
int
sys_write(void)
{
    struct file *f;
    int n;
    char *p;

    if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
        return -1;

    if(!(f->ip->mode & 0x2)){
        cprintf("Operation write failed\n");
        cprintf("mode = %d\n", f->ip->mode);
        return -1;
    }

    return filewrite(f, p, n);
}
```

Figure 19: `sysfile.c` - Checking the permissions before performing read, write, execute operations.

Similarly for `sys_read` and `sys_exec`